

Stefan Indic

[Website](https://stefanind.github.io): stefanind.github.io | [GitHub](https://github.com/stefanind): stefanind | [LinkedIn](https://www.linkedin.com/in/stefanindic): in/stefanindic | stefanindic@gmail.com | Seattle, WA

Education

Master of Science, Data Science Seattle University Seattle, WA	Jun 2026
Bachelor of Arts, Philosophy with Distinction University of British Columbia Vancouver, BC	May 2022
- Relevant Coursework: Calculus III, Linear Algebra, Statistics, and Philosophy of Physics & Mathematics. - Thesis: The Transition from Euclidean to Minkowski Spacetime and the Invariance of the Spacetime Interval.	

Technical Skills

Languages & Tools: Python, SQL, Git, Linux, Docker, FastAPI.
Data Science: Pandas, NumPy, SciPy, Matplotlib, Seaborn, Altair.
Machine Learning: PyTorch, scikit-learn, RAG, NLP, MLflow, MLOps, Prefect.
Cloud & Big Data: AWS, Spark, Snowflake, Hadoop, Hive.

Experience

Data Scientist (Seattle University Capstone) Tacoma Power Seattle, WA	Jan 2026 - June 2026
Architected an end-to-end ML pipeline for physical infrastructure data, applying XGBoost and SHAP to identify the highest-value drivers of electrical pole replacement to help Tacoma Power prioritize inspection and maintenance investments. Collaborated with inspectors, engineers, and senior leadership to validate model assumptions, align outputs with engineering constraints, and translate findings into actionable infrastructure investment recommendations. Preprocessed and standardized multi-year inspection data from Excel and SharePoint into a reusable schema, enabling predictive modeling, repeated analysis, and dashboard reporting at scale. Engineered domain-specific features across pole condition, shell thickness, decay, bird damage, fiber strength, and capacity ratio, assessing multicollinearity via correlation matrix and VIF before modeling. Tuned XGBoost through cross-validation and ranked replacement drivers by SHAP value importance to explain which factors most influenced pole rejection decisions. Presented findings to stakeholders, revealing that 43% of pole rejections stemmed from pole top deterioration caused by decay and bird damage, directly informing maintenance targeting and investment priorities.	
Data Analyst - AI Evaluation DataAnnotation Remote	Feb 2023 - Sep 2024
Evaluated frontier LLM response datasets across coding, statistics, machine learning, reasoning, and agentic tool usage, applying statistical significance testing to compare model outputs and clustering analysis with NLP techniques to identify systematic failure patterns. Designed and audited prompts for RLHF workflows in Python, ML, data structures, and statistics, contributing structured feedback to model training pipelines at scale.	

Projects

stefanGPT https://github.com/stefanind/stefanGPT Try it on my website!	Apr 2026 - Present
Built a production AI twin — fine-tuned Qwen2.5-7B-Instruct via QLoRA + RAG — that lets recruiters assess my technical depth, judgment, and communication style beyond what a resume conveys. Designed the knowledge layer using a curated Q&A fine-tuning dataset, SentenceTransformers embeddings, and FAISS retrieval to ground responses in real personal and professional context. Built deployment and evaluation tooling with Docker, RunPod Serverless, Cloudflare Workers, and custom eval scripts to test answer quality, factual grounding, usefulness, and resemblance to my communication style.	
ML Deployment System https://github.com/stefanind/beginner-ML-deployment	Mar 2026
Built a production-oriented ML deployment system demonstrating FastAPI inference, MLflow experiment tracking, Docker containerization, Prometheus monitoring, Locust load testing, Prefect orchestration, and GitHub Actions CI to show beginners the full path from model training to deployable service.	
Optimization Priors https://github.com/stefanind/research-optimization-priors	Feb 2026 - Present
Designed controlled PyTorch experiments on small transformer models to study how data-derived priors, teacher initialization, hidden-state geometry, and logit distillation affect learning efficiency and optimization dynamics.	