

Stefan Indic

[Website](https://stefanind.github.io): stefanind.github.io | [LinkedIn](https://www.linkedin.com/in/stefanindic): in/stefanindic | (206) 379-2120 | stefanindic@gmail.com

Education

Master of Science, Data Science Seattle University Seattle, WA	Jun 2026
Bachelor of Arts, Philosophy with Distinction University of British Columbia Vancouver, BC	May 2023
Relevant Coursework: Calculus III, Linear Algebra, Statistics, and Philosophy of Physics & Mathematics.	

Technical Skills

Languages & Tools: Python, SQL, Excel, R, Git, Linux, Docker.
Data Science: Pandas, NumPy, SciPy, Matplotlib, Seaborn, A/B testing, Probability & Statistics.
Machine Learning: PyTorch, scikit-learn, NLP, MLflow.
Cloud & Big Data: AWS (EC2, S3), Snowflake, Hadoop, Spark.

Experience

Data Scientist (Seattle University Capstone) Tacoma Power Seattle, WA	Jan 2026 - Present
- Architected the first end-to-end reusable and automated Python ML + analytics reporting pipeline to identify the top drivers of utility pole rejections from historical and future incoming inspection data. - Incorporated LangGraph to orchestrate Excel → SharePoint → Snowflake → Tableau ingestion, validation, auditing, feature engineering, modeling, SHAP explainability, and reporting, allowing a nontechnical user to fix ~80% of pipeline issues by talking with an AI agent. - Collaborated with onsite inspectors, engineers, managers, and senior stakeholders to validate inspection data assumptions, align model outputs with field and engineering workflows, and ensure dashboards answered senior management and C-suite reporting needs. - Defined robust config-driven cleaning and standardization workflows using Pandas, JSON policies, mapping csv files, fuzzy column matching, row filters, audit logs, validation contracts, and unit tests to make current and future messy inspection data reproducible and model-ready. - Implemented NLP-based inspection comment parsing with spaCy and regex rules across 33 comment categories, extracting decay, damage, hardware issues, access constraints, and other risk indicators from unstructured field notes. - Generated interpretable model artifacts with XGBoost + SHAP to identify rejection drivers, and prediction-level explanations, revealing pole top deterioration from decay and bird damage as the leading rejection driver behind 43% of regional rejections.	

Data Analyst DataAnnotation Remote, CA	Feb 2023 - Sep 2024
- Analyzed 10+ frontier LLM response datasets across 7 task domains using NLP and statistical analysis, producing 5+ reports that identified chatbot performance patterns, model weaknesses, and findings such as Model A performing 13% worse than Model B across selected evaluations. - Designed 1,000+ Python, machine learning, software engineering, and statistics prompts for RLHF model evaluation workflows, supporting higher-quality feedback data for frontier LLM training and assessment. - Audited and validated other annotators' work for correctness, instruction adherence, reasoning quality, and task alignment, finding that ~20% of workers missed key task requirements on average.	

Projects

General Purpose Transformer https://github.com/stefanind/GPT	Aug 2025 - Present
- Rebuilt GPT-2 Small 124M from scratch in PyTorch to learn how LLMs work and extended the architecture with modern components including SwiGLU, RMSNorm, KV Cache, GQA/MQA, and Mixture of Experts. - Achieved 0.1770 lower validation loss with only a 0.0011 HellaSwag accuracy delta vs. OpenAI GPT-2 Small, demonstrating stronger performance while maintaining comparable benchmark accuracy. - Improved training throughput by 5× and reached 97% GPU utilization on A100s by profiling bottlenecks with PyTorch Profiler and optimizing bfloat16, TF32, FlashAttention, fused AdamW, batch sizing, and kernel utilization.	
Optimal Cancellation Penalty For Delivery Drivers	Dec 2024
- Ran A/B/n testing of \$0, \$10, and \$20 cancellation penalties for uncompleted deliveries, identifying a significant result of a \$1.5M profit increase from \$5.5M to \$7.0M under the highest penalty scenario. \$10 penalty recommended because it reduced cancellations and improved profit while avoiding the driver engagement risks seen with the \$20 penalty, which could lead to about \$70K in missed revenue over six weeks.	