

Stefan Indic

[Website](https://stefanind.github.io): stefanind.github.io | [GitHub](https://github.com/stefanind): stefanind | [LinkedIn](https://www.linkedin.com/in/stefanindic): in/stefanindic | stefanindic@gmail.com | Seattle, WA

Education

Master of Science, Data Science

Jun 2026

Seattle University | Seattle, WA

- Relevant Coursework: ML I & II, Advanced ML, Deep Learning, Statistical Inference, and A/B testing.

Bachelor of Arts, Philosophy with Distinction

May 2022

University of British Columbia | Vancouver, BC

- Relevant Coursework: Calculus III, Linear Algebra, Statistics, and Philosophy of Physics & Mathematics.

Technical Skills

Languages & Tools: Python, SQL, Git, Linux, Docker, FastAPI.

Data Science: Pandas, NumPy, SciPy, Matplotlib, Seaborn, Altair.

Machine Learning: PyTorch, scikit-learn, RAG, NLP, MLflow, MLOps, Prefect.

Cloud & Big Data: AWS, Spark, Snowflake, Hadoop, Hive.

Experience

Data Scientist (Seattle University Capstone)

Tacoma Power | Seattle, WA

Jan 2026 - Present

- Architected the **first end-to-end reusable and automated Python ML analytics pipeline** for Tacoma Power, **identifying the top causes of electrical pole replacement** across historical and incoming inspection data to help prioritize infrastructure investments and prevent future pole failures.
- Built the pipeline as a 10-stage validated workflow** spanning automated SharePoint ingestion, Excel parsing, data cleaning, schema standardization, NLP comment parsing, feature engineering, XGBoost modeling, SHAP explainability, Snowflake publishing, and Streamlit reporting.
- Validated data cleaning, standardization, and reporting assumptions by **collaborating with onsite inspectors, engineers, and senior management** to ensure the schema reflected field and engineering conditions while model outputs aligned with dashboard reporting needs and operational priorities.
- Revealed that 43% of pole rejections** across the region stem from pole top deterioration caused by decay and bird damage.
- Presented key findings to Tacoma Power stakeholders** through a PowerPoint presentation, translating model insights into infrastructure investment recommendations.

Data Analyst

DataAnnotation | Remote

Feb 2023 - Sep 2024

- Analyzed 10+ frontier LLM response datasets** across 7 task domains using NLP and statistical analysis, **producing 5+ reports that identified chatbot performance** patterns, model weaknesses, and findings such as Model A performed 13% worse than Model B across hyperspecific instruction tasks.
- Designed and annotated 1,000+ Python, machine learning, software engineering, and statistics prompts/tasks for Reinforcement Learning from Human Feedback finetuning.
- Audited and validated other annotators' work** for correctness, instruction adherence, reasoning quality, and task alignment, finding that **~20% of workers missed key task requirements on average**.

Projects

[stefanGPT](https://github.com/stefanind/stefanGPT) | <https://github.com/stefanind/stefanGPT> | Try it on my website!

Apr 2026 - Present

- Created a personal AI system that finetunes Qwen2.5-7B-Instruct with QLoRA + RAG** to model my reasoning style, communication patterns, technical background, and professional judgment, **allowing recruiters to evaluate my fit beyond a static resume**.
- Architected an end-to-end supervised finetuning pipeline that converts my personal Q&A knowledge files into validated chat format JSONL datasets, trains PEFT/LoRA adapters, and stores versioned model artifacts for evaluation and deployment.
- Developed a RAG indexing pipeline using SentenceTransformers and FAISS** to embed curated personal and professional knowledge documents into retrievable chunks with source metadata and reproducible manifests.
- Built deployment and evaluation tooling with Docker, RunPod Serverless, Cloudflare Workers, and custom eval scripts** to test answer quality, factual grounding, usefulness, and resemblance to my communication style.

[ML Deployment System](https://github.com/stefanind/beginner-ML-deployment) | <https://github.com/stefanind/beginner-ML-deployment>

Mar 2026

- Built an ML deployment platform that teaches beginners how to move beyond model training into production** through FastAPI inference, MLflow tracking, Docker containerization, Prometheus monitoring, Locust load testing, Prefect orchestration, and GitHub Actions CI.